



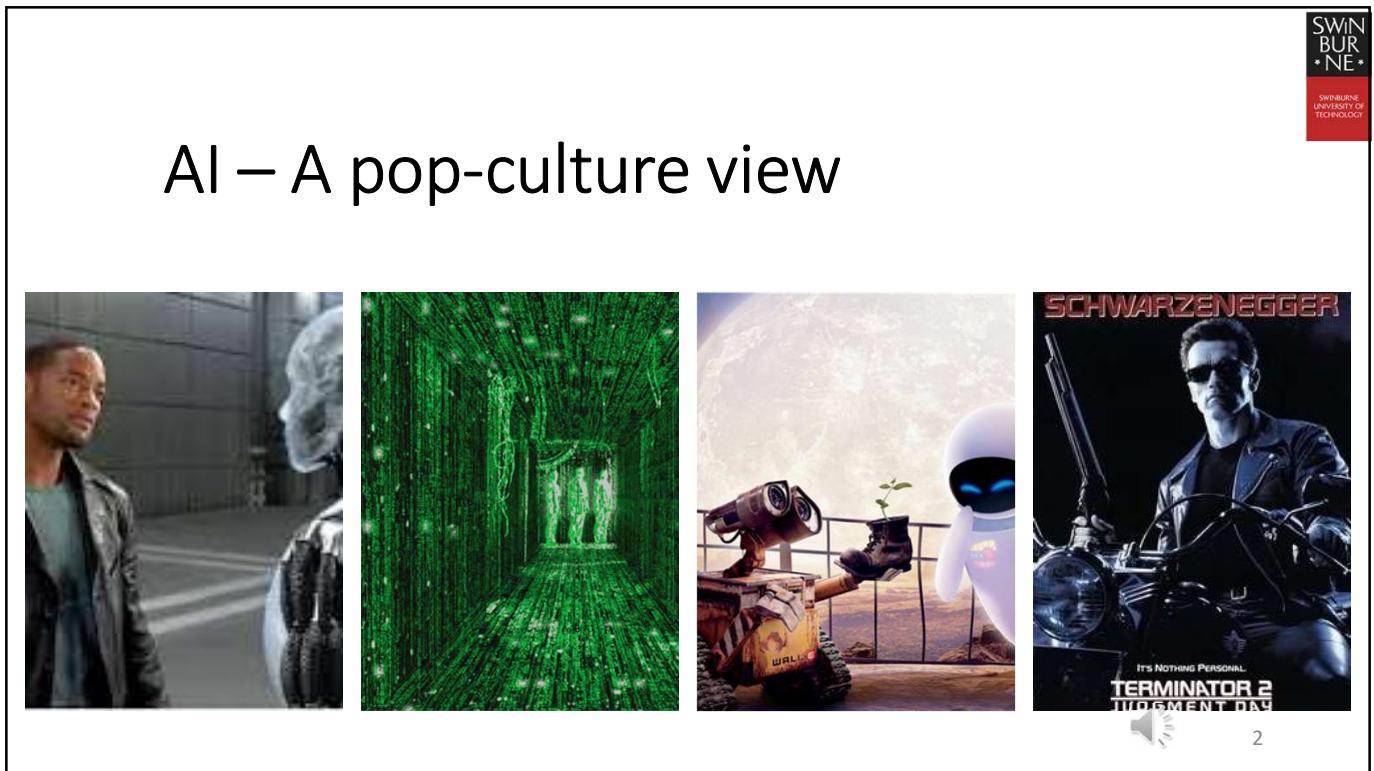
SWIN
BUR
NE

SWINBURNE
UNIVERSITY OF
TECHNOLOGY

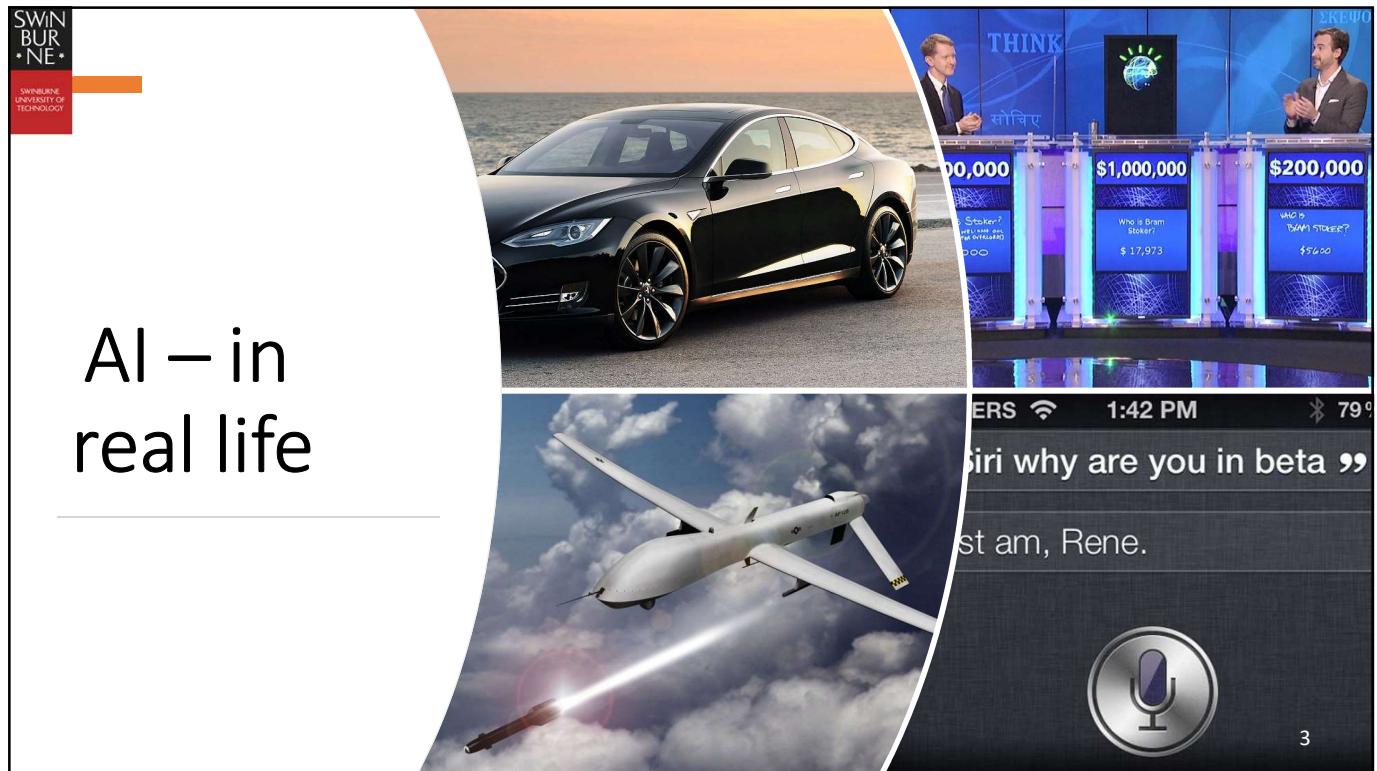
COS30019 – Introduction to AI: Week 1

Introduction

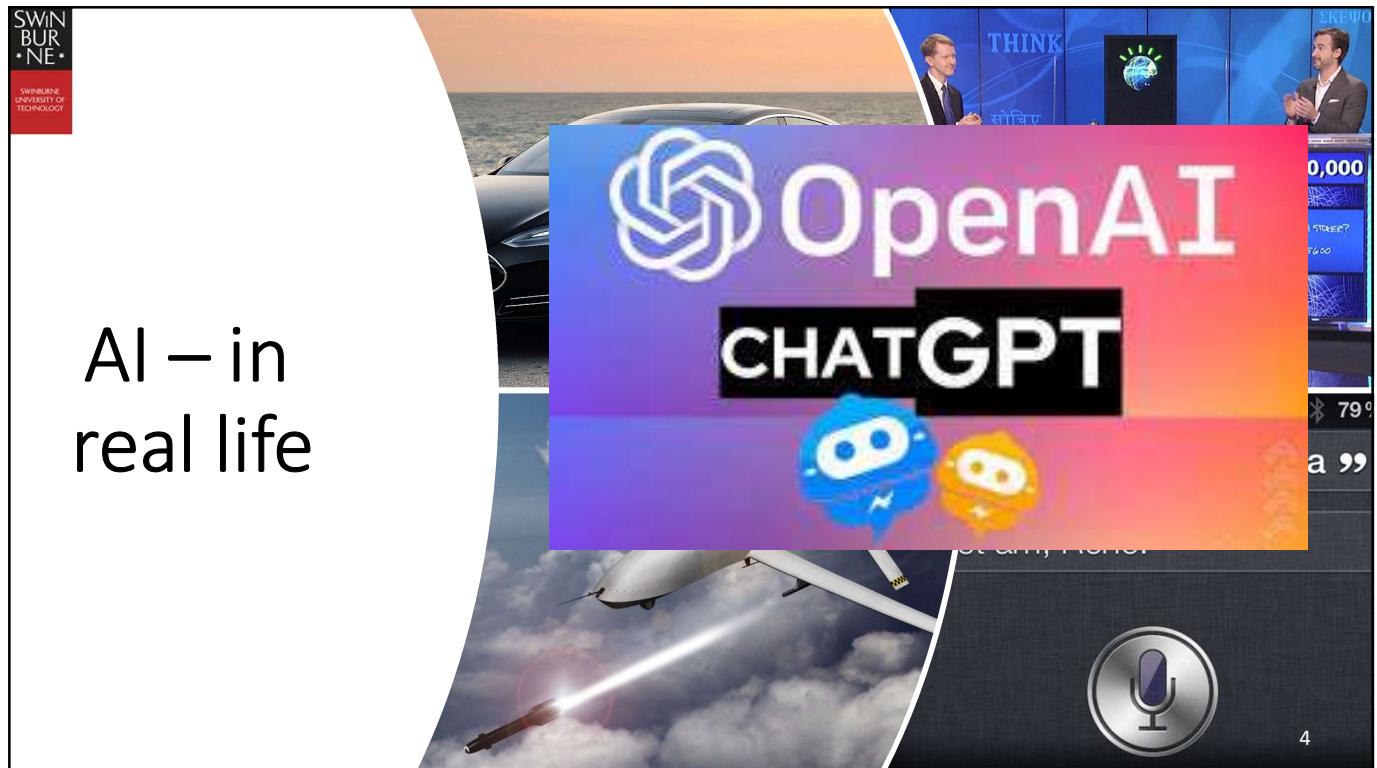
1



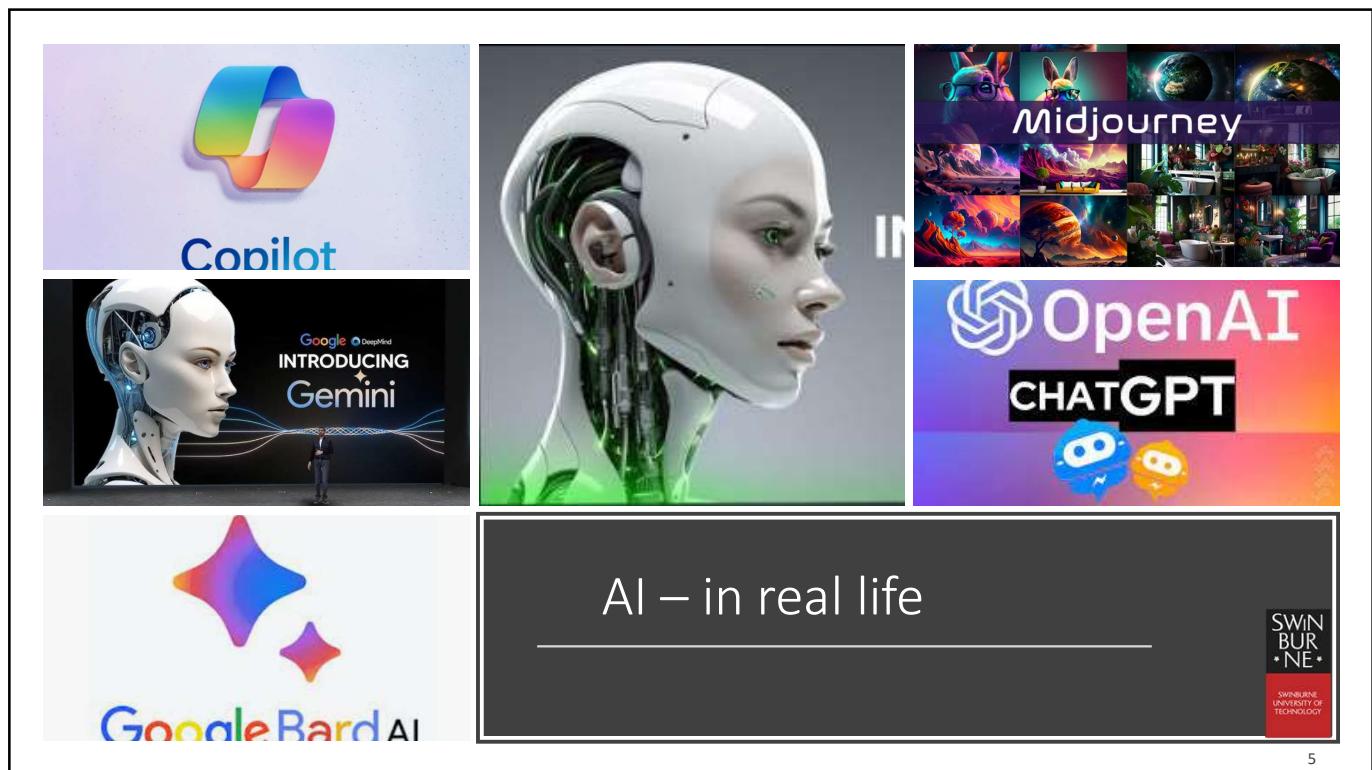
2



3



4



5

What is AI? (Informally)

- Different people can define AI differently
- Most popular definitions:
 - “Artificial intelligence is a constellation of many different technologies working together to enable machines to sense, comprehend, act, and learn with **human-like** levels of intelligence.”
(Accenture)
 - Artificial intelligence is the **simulation of human intelligence** processes by machines, especially computer systems/AI refers to systems or machines that **mimic human intelligence** to perform tasks and can iteratively improve themselves based on the information they collect.
(SAS/Oracle)
 - AI refers to systems that acts **rationally** (aka. Intelligent agents): any system that perceives its environment and takes actions that maximize its chance of achieving its goals
(AI textbooks)

6

Why study AI (from an academic perspective)

- It provides the core knowledge of computer science
- You'll learn to analyse problems and learn about techniques/algorithms to solve real-world problems
- It paves the way to understanding various sorts of intelligence (in both humans and machines)
- It is also fun (and different to most other subjects)

7

Why study AI? (from a career perspective)

- The market size in the AI market is projected to reach US\$305.90bn in 2024. The market size is expected to show an annual growth rate (CAGR 2024-2030) of 15.83%, resulting in a market volume of US\$738.80bn by 2030.
(Statista)
- The total contribution of AI to the global economy is expected to hit \$15.7 trillion by 2030.
(PwC Global)
- AI will help boost the GDP of local economies, with China expected to record the greatest gains of 26% by 2030.
(PwC Global)
- The most in-demand AI job of 2023 can pay over US\$200,000 and offers remote opportunities
(CNBC, Nov 2023)
- Searches for generative AI jobs on Indeed have increased almost 4,000% in the last year, and openings for generative AI jobs are up 306% over the same period.
(CNBC, Nov 2023)

8



What is Artificial Intelligence

- Different definitions due to different criteria
 - Two dimensions:
 - Thought processes/reasoning vs. behavior/action
 - Success according to human standards vs. success according to an ideal concept of intelligence: rationality.

Systems that think like humans	Systems that think rationally
Systems that act like humans	Systems that act rationally

- Each definition falls under a **PARADIGM** in which AI can be built

Systems that act rationally

- Rational behavior: “doing the right thing”
 - The “Right thing” is the course of action that is expected to *maximize goal achievement given the available information*.
- Can include thinking, yet in service of rational action.
 - Action without thinking: e.g. reflexes.

Systems that act rationally (aka. **Intelligent agents**)

Systems that act rationally

- Two advantages over other approaches:
 - More general than law of thoughts approach
 - More amenable to scientific development (e.g., not constrained by how humans think/act).
- Yet, rationality is only applicable in *ideal* environments.
- Moreover, rationality may not be a good model of reality.

11

Foundations of AI

- Different fields have contributed to AI in the form of ideas, viewpoints and techniques.
 - *Philosophy*: Logic, reasoning, mind as a physical system, foundations of learning, language and rationality.
 - *Mathematics*: Formal representation and proof algorithms, computation, (un)decidability, (in)tractability, probability.
 - *Psychology*: adaptation, phenomena of perception and motor control.
 - *Economics*: formal theory of rational decisions, game theory.
 - *Linguistics*: knowledge representation, grammar.
 - *Neuroscience*: physical substrate for mental activities.
 - *Control theory*: homeostatic systems, stability, optimal agent design.

12

History of AI

Early computing

First Golden Age
1956 - 1974

AI Winter

Quiet Progress

Second Golden Age



Dartmouth Conference, 1956

"Every aspect of learning or any other feature of intelligence can be so precisely described that a machine can be made to simulate it"

13

History of AI

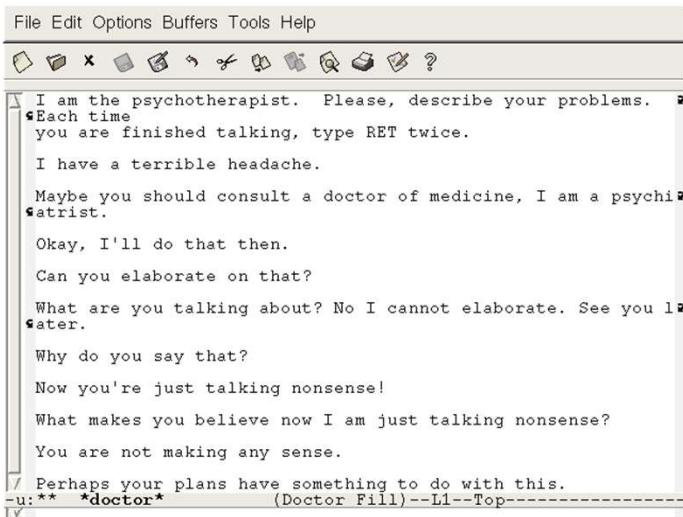
Early computing

First Golden Age
1956 - 1974

AI Winter

Quiet Progress

Second Golden Age



ELIZA, 1964

14

History of AI

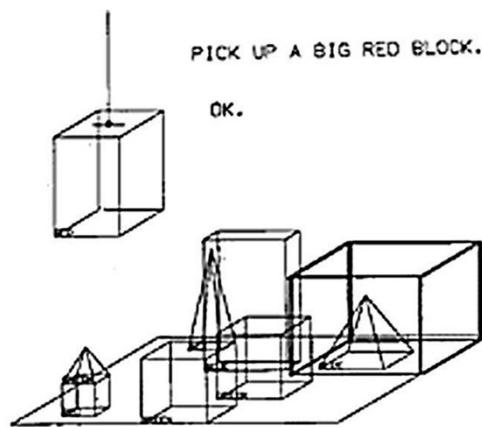
Early computing

First Golden Age
1956 - 1974

AI Winter

Quiet Progress

Second Golden Age



SHRDLU, 1968

15

History of AI

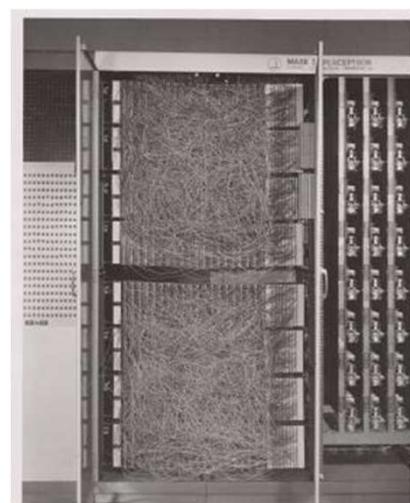
Early computing

First Golden Age
1956 - 1974

AI Winter

Quiet Progress

Second Golden Age



Perceptron, 1957

Attendance word: CHECK SLIDES

16

History of AI

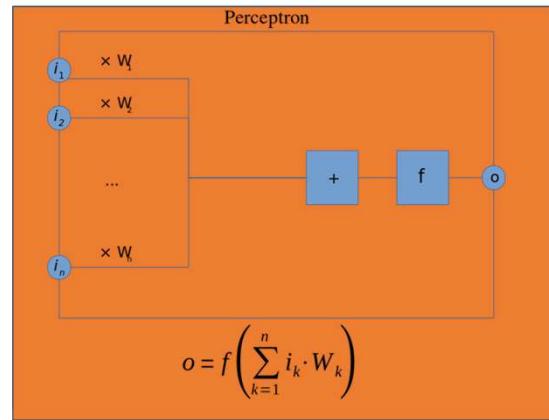
Early computing

First Golden Age
1956 - 1974

AI Winter

Quiet Progress

Second Golden Age



Perceptron, 1957

17

History of AI

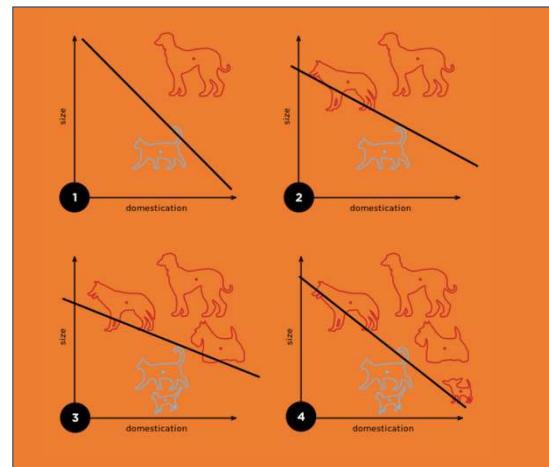
Early computing

First Golden Age
1956 - 1974

AI Winter

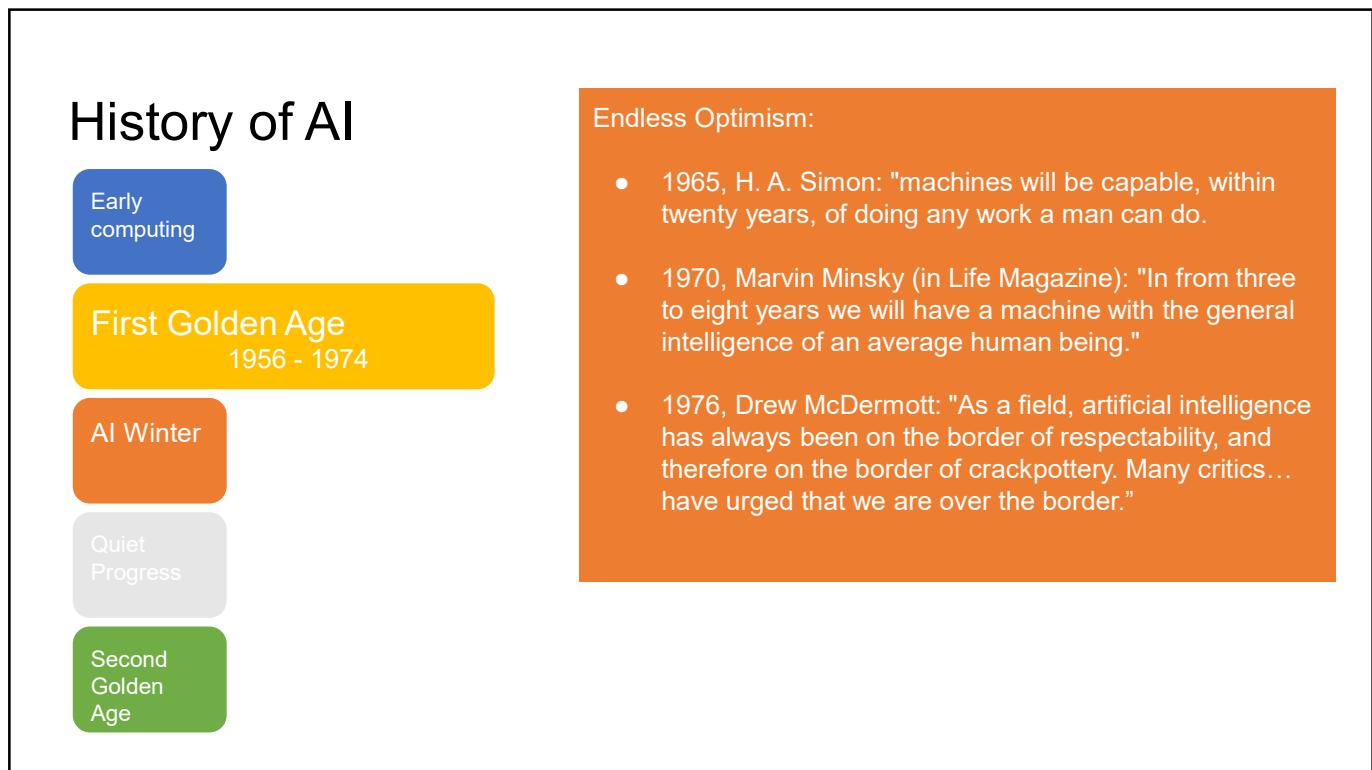
Quiet Progress

Second Golden Age

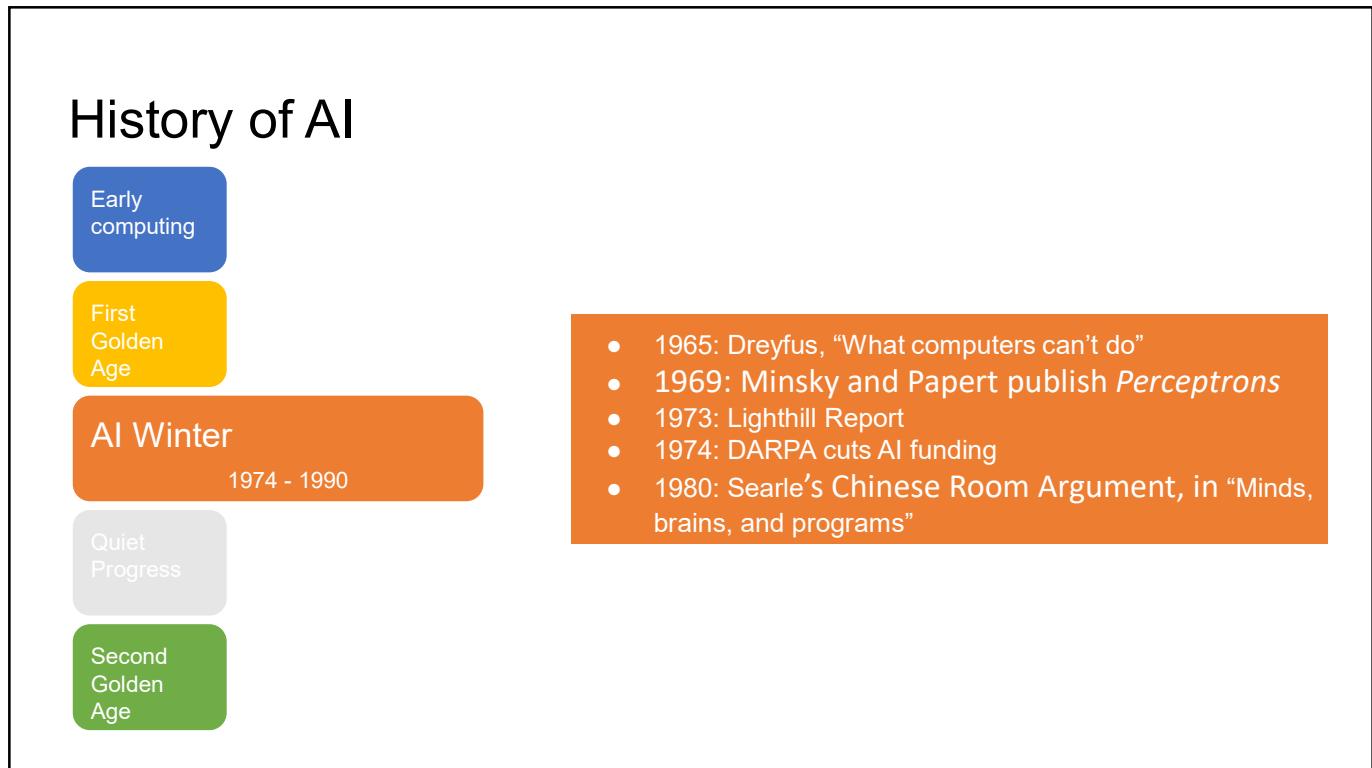


Perceptron, 1957

18



19



20

History of AI

Early computing

First Golden Age

AI Winter

Quiet Progress

1990 - 2012

Second Golden Age



[Kasparov vs Deep Blue](#)

May 11, 1997

(the day computers took over the world)

21

History of AI

Early computing

First Golden Age

AI Winter

Quiet Progress

1990 - 2012

Second Golden Age



[Watson vs Jeopardy](#) 2011

Smartest Machine on Earth

(NOVA documentary)

22

History of AI

Early computing

First Golden Age

AI Winter

Quiet Progress

1990 - 2012

Second Golden Age



Roomba, 2002, based on Rodney Brooks' subsumption architecture

23

History of AI

Early computing

First Golden Age

AI Winter

Quiet Progress

1990 - 2012

Second Golden Age



DARPA Grand Challenge, 2005
[Great Robot Race](#) (NOVA documentary)

24

History of AI

Early computing

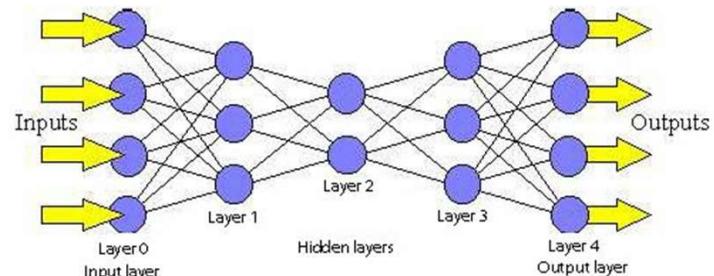
First Golden Age

AI Winter

Quiet Progress

1990 - 2012

Second Golden Age



Back-propagation with hidden layers, 1986

25

History of AI

Early computing

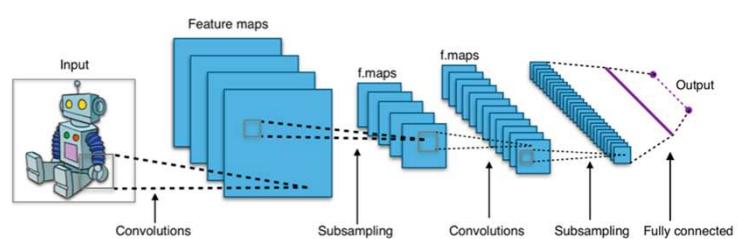
First Golden Age

AI Winter

Quiet Progress

1990 - 2012

Second Golden Age



Yann LeCun, 1998
Convolutional Neural Networks

26

History of AI

Early computing

First Golden Age

AI Winter

Quiet Progress

1990 - 2012

Second Golden Age



Geoffrey Hinton talks
@Google, 2007

27

History of AI

Early computing

First Golden Age

AI Winter

Quiet Progress

Second Golden Age
2012 - Present



Geoffrey Hinton beats
ImageNet, 2012

28

History of AI

Early computing

First Golden Age

AI Winter

Quiet Progress

Second Golden Age
2012 - Present

- 2012: Hinton uses Deep ConvNets to solve ImageNet classification
- 2013: Hinton hired by Google, uses NN on mobile
 - Boston Dynamics acquired by Google
 - Yann LeCun heads FAIR for Facebook
- 2014: Google acquires DeepMind
- 2015: Musk donates \$10m to Future of Life to study existential risk for artificial intelligence
 - Autonomous vehicle testing begins
- 2016: AlphaGo beats Lee Sedol
- 2017: AlphaZero beats Stockfish
- 2022: OpenAI releases ChatGPT
- ... : OpenAI's GPT4o, Google's Bard AI, Google's Gemini, Meta's Llama, SORA, etc.

29

Different types of current AI systems

- **Rule-based** systems (e.g., expert systems, Mars Rover AI planner)
- **Machine learning (ML)-based** systems (e.g., IBM Watson, data analytics systems used by supermarkets, insurance and banking sector, telcos, etc.)
- **Deep learning (DL)-based** systems (e.g., those used in autonomous cars, Alpha Go, ChatGPT, etc.)

30

Some other terminologies



- Strong AI **vs** Weak AI
 - Weak AI:
 - Machines that can be made to act *as if* they were intelligent.
 - **Strong AI:**
 - **Machines that act intelligently with real, conscious minds.**

- Narrow AI **vs** Artificial General Intelligence (AGI)
 - Narrow AI:
 - Machine that is focused on one narrow (intellectual) task.
 - AGI:
 - Machine with the ability to apply intelligence to any problem, rather than just one specific problem.

Ethics and AI

- **Ethical considerations** in AI:
 - Bias
 - Privacy
 - job displacement
- **Examples** of ethical issues in AI
 - facial recognition
 - algorithmic hiring
- **Potential solutions to ethical challenges**
 - European Commission (EC)'s AI Act (AIA)
 - Responsible AI
 - Transparent & Auditable AI

Future of AI

- In your hand
 - You'll create it!
- Topic for discussion in class (tutorial)

33

Summary

- Different people think of AI differently.
- Two important questions to ask are:
- Are you concerned with thinking or behavior?
- Do you want to model humans or work from an ideal standard?
- In this course, we adopt the view that intelligence is concerned mainly with rational action.
- Ideally, an intelligent agent takes the best possible action in a situation. We will study the problem of building agents that are intelligent in this sense.



34



35

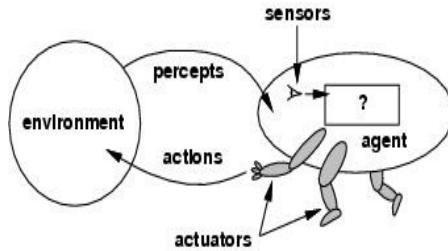
Outline

- Agents and environments.
 - The vacuum-cleaner world
- The concept of rational behavior.
- Environments.
- Agent structure.

36

Agents and environments

- Agents include human, robots, softbots, thermostats, etc.
- The *agent function* maps percept sequence to actions



- An agent can perceive its own actions, but not always its effects.

$$f : P^* \rightarrow A$$

37

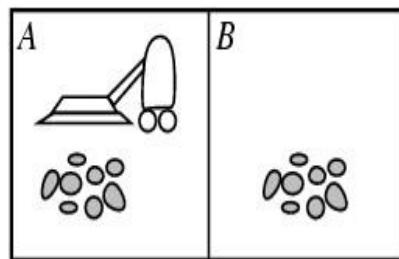
To refresh your memory

- This is a mathematical function (square):
 - $f: \mathbb{R} \rightarrow \mathbb{R}$
 - $f(x) = x^2$
- To explain this square function to primary school students, teacher use the following table:

x	x^2
1	1
2	4
3	9
4	16
5	25
6	36
...	...

38

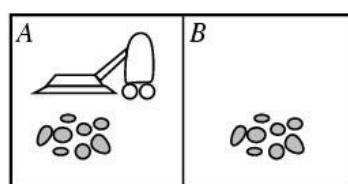
The vacuum-cleaner world – An example



- **Environment:** squares A and B
- **Percepts:** [location and content] e.g. [A, Dirty]
- **Actions:** left, right, suck, and no-op

39

The vacuum-cleaner world – Agent function



Percept sequence	Action
[A, Clean]	Right
[A, Dirty]	Suck
[B, Clean]	Left
[B, Dirty]	Suck
[A, Clean],[A, Clean]	Right
[A, Clean], [A, Dirty]	Suck
...	...

40



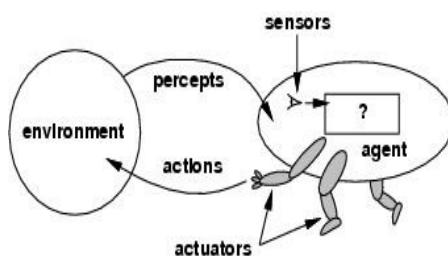
Hang on a second

- Why does this sound familiar?
- For Generative AI systems (e.g., ChatGPT, Gemini, DeepSeek, ...):
 - The prompts ~ percept sequences
 - AI system outputs ~ Agent's actions
- **Agentic AI**

41

Agents and environments

- The *agent function* will internally be represented by the *agent program*.
- The agent program runs on the physical *architecture* to produce f .



42

The concept of rationality

- A **rational agent** is one that does the right thing.
 - Every entry in the table is filled out correctly.
- What is the right thing?
 - Approximation: the most *successful* agent.
 - *Measure of success?*
- Performance measure should be **objective**
 - E.g. the amount of dirt cleaned within a certain time.
 - E.g. **how clean the floor is.**
 - ...
- **Performance measure according to what is wanted in the environment instead of how the agents should behave.**

43

Rationality

- What is rational at a given time depends on four things:
 - Performance measure,
 - Prior environment knowledge,
 - Actions,
 - Percept sequence to date (sensors).
- DEF: *A rational agent chooses whichever action that maximizes the expected value of the performance measure given the percept sequence to date and prior environment knowledge.*

44

Rationality

- The proposed definition requires:
 - Information gathering/exploration
 - To maximize future rewards
 - Learn from percepts
 - Extending prior knowledge
 - Agent autonomy
 - Compensate for incorrect prior knowledge

45

Is the vacuum cleaner agent rational?

- Depend!
- For example, it's rational under the following assumptions:
 - Performance measure: 1 point for each clean square over 'lifetime' of 1000 steps
 - 'geography' known but dirt distribution, initial position of agent not known
 - Clean squares stay clean, sucking cleans squares
 - Left and Right don't take agent outside environment
 - Available actions: Left, Right, Suck, NoOp
 - Agent knows where it is and whether that location contains dirt
- Let's revisit this during the tutorial

46

Environments

- To design a rational agent, we must specify its **task environment**
- PEAS description of the task environment:
 - Performance
 - Environment
 - Actuators
 - Sensors

47

Environments

- E.g. Fully automated taxi:
 - PEAS description of the environment:
 - Performance
 - Safety, destination, profits, legality, comfort
 - Environment
 - Streets/freeways, other traffic, pedestrians, weather, ...
 - Actuators
 - Steering, accelerating, brake, horn, speaker/display,...
 - Sensors
 - Video, sonar, speedometer, engine sensors, keyboard, GPS, ...

48

Environment types

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable?? Deterministic?? Episodic?? Static?? Discrete?? Single-agent??				

49

The game of backgammon



50

Environment types

Fully vs. partially observable: an environment is full observable when the sensors can detect all aspects that are *relevant* to the choice of action.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable?? Deterministic?? Episodic?? Static?? Discrete?? Single-agent??				

51

Environment types

Fully vs. partially observable: an environment is full observable when the sensors can detect all aspects that are *relevant* to the choice of action.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable?? Deterministic?? Episodic?? Static?? Discrete?? Single-agent??	FULL	FULL	FULL	PARTIAL

52

Environment types

Deterministic vs. stochastic: if the next environment state is completely determined by the current state the executed action then the environment is deterministic.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable?? Deterministic?? Episodic?? Static?? Discrete?? Single-agent??	FULL	FULL	FULL	PARTIAL

53

Environment types

Deterministic vs. stochastic: if the next environment state is completely determined by the current state the executed action then the environment is deterministic.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable?? Deterministic?? Episodic?? Static?? Discrete?? Single-agent??	FULL YES	FULL NO	FULL YES	PARTIAL NO

54

Environment types

Episodic vs. sequential: In an episodic environment the agent's experience can be divided into atomic steps where the agents perceives and then performs A single action. The choice of action depends only on the episode itself

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??				
Static??				
Discrete??				
Single-agent??				

55

Environment types

Episodic vs. sequential: In an episodic environment the agent's experience can be divided into atomic steps where the agents perceives and then performs A single action. The choice of action depends only on the episode itself

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??				
Discrete??				
Single-agent??				

56

Environment types

Static vs. dynamic: If the environment can change while the agent is choosing an action, the environment is dynamic. Semi-dynamic if the agent's performance changes even when the environment remains the same.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??				
Discrete??				
Single-agent??				

57

Environment types

Static vs. dynamic: If the environment can change while the agent is choosing an action, the environment is dynamic. Semi-dynamic if the agent's performance changes even when the environment remains the same.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??				
Single-agent??				

58

Environment types

Discrete vs. continuous: This distinction can be applied to the state of the environment, the way time is handled and to the percepts/actions of the agent.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??				
Single-agent??				

Environment types

Discrete vs. continuous: This distinction can be applied to the state of the environment, the way time is handled and to the percepts/actions of the agent.

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??	YES	YES	YES	NO
Single-agent??				

Environment types

Single vs. multi-agent: Does the environment contain other agents who are also maximizing some performance measure that depends on the current agent's actions?

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??	YES	YES	YES	NO
Single-agent??				

61

Environment types

Single vs. multi-agent: Does the environment contain other agents who are also maximizing some performance measure that depends on the current agent's actions?

	Crossword	Backgammon	Chess w/ clock	Taxi
Observable??	FULL	FULL	FULL	PARTIAL
Deterministic??	YES	NO	YES	NO
Episodic??	NO	NO	NO	NO
Static??	YES	YES	SEMI	NO
Discrete??	YES	YES	YES	NO
Single-agent??	YES	NO	NO	NO

62

Environment types

- The simplest environment is
 - Fully observable, deterministic, episodic, static, discrete and single-agent.
- Most real situations are:
 - Partially observable, stochastic, sequential, dynamic, continuous and multi-agent.

63

Agent types

- How does the inside of the agent work?
 - Agent = architecture + program
- All agents have the same skeleton:
 - Input = current percepts
 - Output = action
 - Program= manipulates input to produce output
- Note difference with agent function.

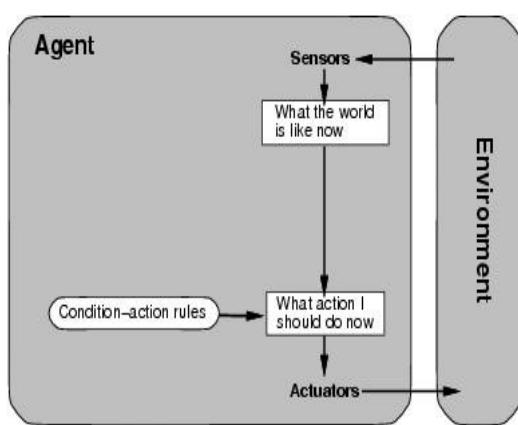
64

Agent types

- Four **basic** kind of agent programs will be discussed:
 - Simple reflex agents
 - Model-based reflex agents
 - Goal-based agents
 - Utility-based agents
- All these can be turned into learning agents.
 - And that gives you **four** additional **advanced agent types**

65

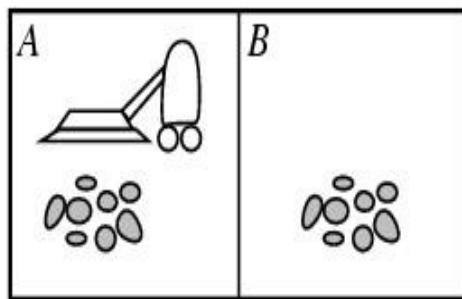
Agent types; simple reflex



- Select action on the basis of *only the current* percept.
 - E.g. the vacuum-agent
- Large reduction in possible percept/action situations(next page).
- Implemented through *condition-action rules*
 - If dirty then suck

66

The vacuum-cleaner world



```
function REFLEX-VACUUM-AGENT ([location, status]) return an action
  if status == Dirty then return Suck
  else if location == A then return Right
  else if location == B then return Left
```

Reduction from 4^T to 4 entries

67

Agent types; simple reflex

function SIMPLE-REFLEX-AGENT(*percept*) returns an action

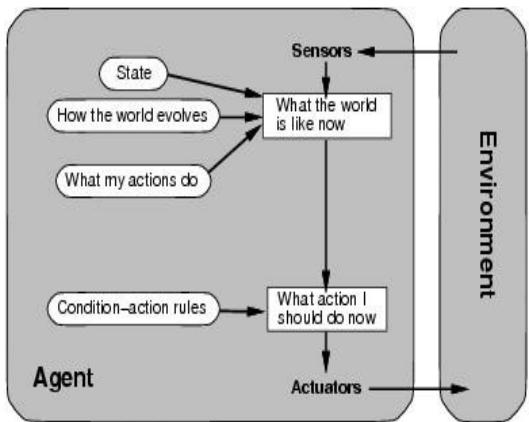
static: *rules*, a set of condition-action rules

```
state  $\leftarrow$  INTERPRET-INPUT(percept)
rule  $\leftarrow$  RULE-MATCH(state, rules)
action  $\leftarrow$  RULE-ACTION[rule]
return action
```

Will only work if the environment is *fully observable* otherwise infinite loops may occur.

68

Agent types; reflex and state



- To tackle *partially observable* environments.
 - Maintain internal state
 - Over time update state using world knowledge
 - How does the world change.
 - How do actions affect world.
- ⇒ *Model of World*

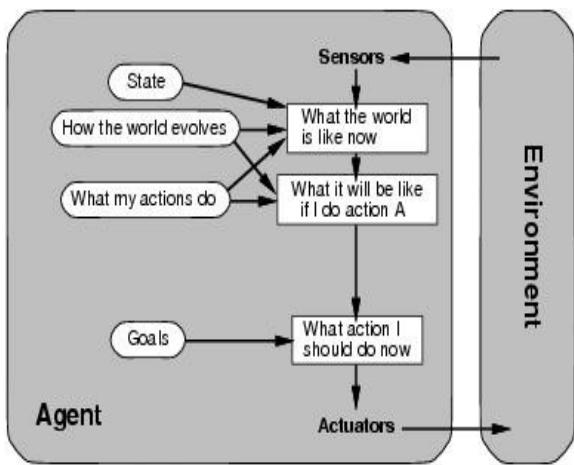
69

Agent types; reflex and state

```
function REFLEX-AGENT-WITH-STATE(percept) returns an action
  static: rules, a set of condition-action rules
    state, a description of the current world state
    actions, the most recent actions.
  state ← UPDATE-STATE(state, actions, percept)
  rule ← RULE-MATCH(state, rules)
  action ← RULE-ACTION[rule]
  return action
```

70

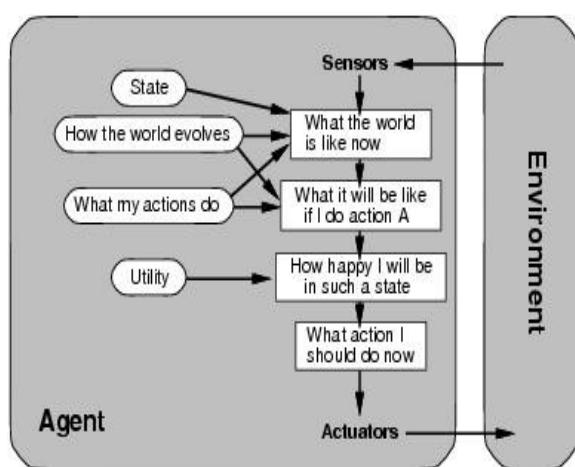
Agent types; goal-based



- The agent needs a goal to know which situations are *desirable*.
 - Things become difficult when long sequences of actions are required to find the goal.
- Typically investigated in **search** and **planning** research.
- Major difference: future is taken into account
- Is more flexible since knowledge is represented explicitly and can be manipulated.

71

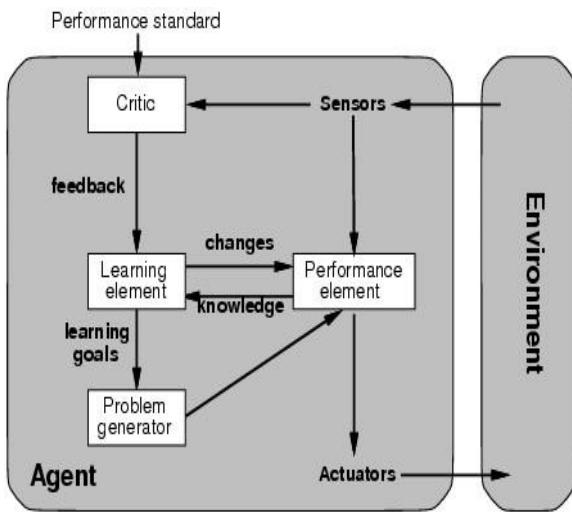
Agent types; utility-based



- Certain goals can be reached in different ways.
 - Some are better, have a higher utility.
- Utility function maps a (sequence of) state(s) onto a real number.
- Improves on goals:
 - Selecting between conflicting goals
 - Select appropriately between several goals based on likelihood of success.

72

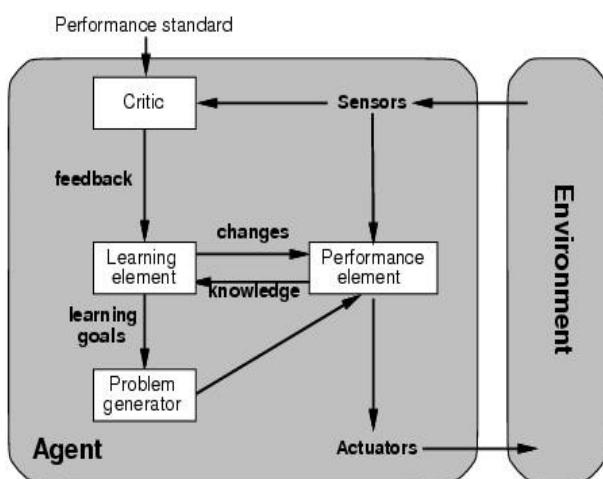
Agent types; learning



- All previous agent-programs describe methods for selecting *actions*.
 - Yet it does not explain the origin of these programs.
 - Learning mechanisms can be used to perform this task.
 - Teach them instead of instructing them.
 - Advantage is the robustness of the program toward initially unknown environments.

73

Agent types; learning



- *Learning element*: introduce improvements in performance element.
 - Critic provides feedback on agents performance based on fixed performance standard.
- *Performance element*: selecting actions based on percepts.
 - Corresponds to the previous agent programs
- *Problem generator*: suggests actions that will lead to new and informative experiences.
 - Exploration vs. exploitation

74

Summary: Agents

- An **agent** perceives and acts in an environment, has an architecture, and is implemented by an agent program.
- Task environment – **PEAS** (Performance, Environment, Actuators, Sensors)
- An **ideal agent** always chooses the action which maximizes its expected performance, given its percept sequence so far.
- An **autonomous learning agent** uses its own experience rather than built-in knowledge of the environment by the designer.
- An **agent program** maps from percept to action and updates internal state.
 - **Reflex agents** respond immediately to percepts.
 - **Goal-based agents** act in order to achieve their goal(s).
 - **Utility-based agents** maximize their own utility function.
- **Representing knowledge** is important for successful agent design.
- The most challenging environments are not fully observable, nondeterministic, dynamic, and continuous

75

References/Readings

1. <https://deeperinsights.com/ai-blog/the-intelligent-agents-of-tomorrow-a-guide-to-llm-powered-agents#:~:text=LLM%2Dpowered%20agents%20function%20as,until%20the%20problem%20is%20resolved>
2. <https://bair.berkeley.edu/blog/2024/02/18/compound-ai-systems/>



76